

```
In [1]: # A dependency of the preprocessing for BERT inputs
!pip install -q tensorflow-text
```

```
4.3MB 2.1MB/s
454.3MB 39kB/s
471kB 37.7MB/s
6.0MB 31.0MB/s
4.0MB 33.2MB/s
4.0MB 22.8MB/s
1.2MB 30.5MB/s
4.9MB 30.8MB/s
```

```
In [2]: # Using AdamW optimizer
!pip install -q tf-models-official==2.4
```

```
1.1MB 2.8MB/s
38.2MB 71kB/s
102kB 10.0MB/s
51kB 6.3MB/s
358kB 34.0MB/s
645kB 30.8MB/s
174kB 42.1MB/s
686kB 34.8MB/s
1.2MB 37.4MB/s
```

Building wheel for py-cpuinfo (setup.py) ... done

Building wheel for sequeval (setup.py) ... done

```
In [4]: import os
import shutil

import tensorflow as tf
import tensorflow_hub as hub
import tensorflow_text as text
from official.nlp import optimization # to create AdamW optimizer

import matplotlib.pyplot as plt

tf.get_logger().setLevel('ERROR')
```

```
In [7]: url = 'https://github.com/ahltraf/point/blob/main/civ_unciv_emails2.tar.gz?raw=true'
dataset = tf.keras.utils.get_file('civ_unciv_emails2.tar.gz', url,
                                untar=True, cache_dir='.',
                                cache_subdir='')
```

Downloading data from https://github.com/ahltraf/point/blob/main/civ_unciv_emails2.ta
r.gz?raw=true

65536/62597 [=====] - 0s 0us/step

```
In [8]: dataset_dir = os.path.join(os.path.dirname(dataset), 'civ_unciv_emails2')
train_dir = os.path.join(dataset_dir, 'train')
```

```
In [9]: AUTOTUNE = tf.data.AUTOTUNE
batch_size = 32
seed = 42

raw_train_ds = tf.keras.preprocessing.text_dataset_from_directory(
    'civ_unciv_emails2/train',
    batch_size=batch_size,
    validation_split=0.2,
    subset='training',
    seed=seed)

class_names = raw_train_ds.class_names
train_ds = raw_train_ds.cache().prefetch(buffer_size=AUTOTUNE)
```

```

val_ds = tf.keras.preprocessing.text_dataset_from_directory(
    'civ_unciv_emails2/train',
    batch_size=batch_size,
    validation_split=0.2,
    subset='validation',
    seed=seed)

val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)

test_ds = tf.keras.preprocessing.text_dataset_from_directory(
    'civ_unciv_emails2/test',
    batch_size=batch_size)

test_ds = test_ds.cache().prefetch(buffer_size=AUTOTUNE)

```

Found 135 files belonging to 2 classes.
 Using 108 files for training.
 Found 135 files belonging to 2 classes.
 Using 27 files for validation.
 Found 33 files belonging to 2 classes.

Looking at and preprocessing emails:

```

In [10]: import re
import nltk
nltk.download('stopwords')
from nltk.corpus import stopwords

```

[nltk_data] Downloading package stopwords to /root/nltk_data...
 [nltk_data] Unzipping corpora/stopwords.zip.

```

In [11]: regex_tokenizer = nltk.RegexpTokenizer("\w+")
def text_preprocessing(content):
    content = str(content)
    content = re.sub("[^a-zA-Z]", " ", content)
    content = content.lower()
    content = content.encode("utf-8", "ignore").decode()
    content = " ".join(regex_tokenizer.tokenize(content))
    for c in content:
        c.replace('\n', ' ')
    words = content.split()
    stops = set(stopwords.words("english"))
    words = [w for w in words if not w in stops]

    return ' '.join(words)

train_2 = train_ds
for text_batch, label_batch in train_2:
    text_batch = text_preprocessing(text_batch)

for text_batch, label_batch in train_2.take(1):
    for i in range(10):
        print(f'Email: {text_batch.numpy()[i]}')
        label = label_batch.numpy()[i]
        print(f'Label: {label} ({class_names[label]})')

```

Email: b"This is indeed guaranteed. For FTRACE use case. If it's being called from FTRACE inrun time, this would mean there were long calls in this module section, which inturn means, get_module_plt() was called at least once for this module and this section. This doesn't hold in general, though. In any case, if you insist, I can try to rework the whole stuff implementing module_finalize().--Best regards,Alexander Sverdlin."

Label: 0 (civil)

Email: b'Are you going to answer any of my remaining questions in a more constructive way?Regards,Markus'

Label: 1 (uncivil)

Email: b'Michal Hocko wrote:Then, I am wondering why we are holding mmap_sem when ca

llingmigrate_pages() in existing code.<http://elixir.free-electrons.com/linux/latest/source/mm/migrate.c#L1576>Sorry, I missed that. If mmap_sem is not needed for migrate_pages(), please ignore this patch.--Best Regards,Yan Zi'

Label: 0 (civil)

Email: b"IMO symlinks are mostly ending in a mess, URLs are never stable. There is a <https://www.kernel.org/doc/html/latest/objects.inv> to handle such requirements. Take a look at *intersphinx* : <http://www.sphinx-doc.org/en/stable/ext/intersphinx.html> to see how it works: Each Sphinx HTML build creates a file named objects.inv that contains a mapping from object names to URIs relative to the HTML set's root. This means articles from external (like lwn articles) has to be recompiled. Not perfect, but a first solution. I really like them, factually valuable comments .. please express your concern so that we have a chance to move on. I think that's a pity.-- Markus --"

Label: 1 (uncivil)

Email: b"it would be very helpful if you cc all involved people on the cover letter instead of just cc'ing your own pile of email addresses. CC'ed now. This is really not helpful. The cover letter and the change logs should contain a summary of that discussion and a proper justification of the proposed change. Just saying 'sysadmins might want to allow' is not useful at all, it's yet another 'I want a pony' thing. I read through the previous thread and there was a clear request to involve security people in to this. Especially those who are deeply involved with hardware side channels. I don't see anyone Cc'ed on the whole series. For the record, I'm not buying the handwavy 'more noise' argument at all. It wants a proper analysis and we need to come up with criteria which PMUs can be exposed at all. All of this wants a proper documentation clearly explaining the risks and scope of these knobs per PMU. Just throwing magic knobs at sysadmins and then saying 'it's their problem to figure it out' is not acceptable. Thanks, \ttglx"

Label: 1 (uncivil)

```

Email: b'Thank you so much for many style, formatting and other issues fixes and als
o for integration of \'check_at_most_once\' patch, it saved me several review iterati
ons. Regarding free of sg in two error paths, you were correct. I fixed it by placin
g several error labels to differentiate each handling. I also noted that reqdata_arr
[b].req was not released properly, this is also fixed. following is a diff of my fix
based on your modifications. (I can send it in a patch format, but it doesn\'t includ
e a fix for Eric Biggers comments)@@ -573,10 +573,9 @@ static void verity_verify_io
(struct dm_verity_io *io)
verity_bv_skip_block(v, io, &io->iter);
continue;
}- reqdata_arr
[b].req = ahash_request_alloc(v->tfm, GFP_NOIO);
if (unlikely(reqdata
_arr[b].req == NULL))-
goto err_memfree;+
ahash_request_set_tfm(reqdata_arr[b].req, v->tfm);
/* +1 for the salt buffer */@@ -586,7 +585,7 @@ static void verity_verify_io(struct
dm_verity_io *io)
GFP_NOIO);
if (!s
g) {
DMERR_LIMIT("%s: kmalloc_array failed", __func__);
goto err_memfree;+
goto err_mem_sg;
}
sg_init_table(sg, num_of_buffs);
// FIXME: if we \'err_memfree\' (or
continue;) below how does this sg get kfree()\'d?@@ -595,7 +594,7 @@ static void ver
ity_verify_io(struct dm_verity_io *io)
reqd
ata_arr[b].want_digest,
&reqdata_arr[b].fec
_io, &is_zero);
if (unlikely(r < 0))-
goto err_
memfree;+
goto err_mem;
if (is_zero) {
/*@@ -605,7 +604,7 @@ static void verity_verify_io(struct dm_verity_io *io)
r = verity_for_bv_block(v, io, &io->iter,
verity_bv_zero);
if (unlikely(r < 0))-
goto err_memfree;+
goto err_mem;
verity_cb_complete(iodata, r);
continue;
}@@ -
644,7 +643,11 @@ static void verity_verify_io(struct dm_verity_io *io)
}
return; -err_memfree: +err_mem: +
kfree(sg); +err_mem_sg: +
ahash_request_fre
e(reqdata_arr[b].req); +err_mem_req:
/*
* reduce expected requests by
the number of unsent
* requests, -1 accounting for the current block
atomic_sub(blocks - b - 1, &iodata->expected_reqs);
verity_cb_complete(iodat
a, -EIO); I took your modifications and working upon it.'

```

Label: 0 (civil)

Email: b'Hi!This is better than my proposal. Thanks!\t\t\t\t\tPavel--(english)
http://www.livejournal.com/~pavelmachek(cesky, pictures) http://atrey.karlin.mff.cun
i.cz/~pavel/picture/horses/blog.html'

Label: 0 (civil)

```
Email: b"I can't take patches without any changelog text at all :("

```

Label: 0 (civil)

Email: b"Ah only if google could simply answer all our questions!It's not like there

is or isn't a security risk and that you can say that it is or it isn't in a global way. Essentially these are channels of information. The channels always exist in form of timing variances for any shared resource (like shared caches or shared memory/I/O/interconnect bandwidth) that can be measured. Perfmon counters make the channels generally less noisy, but they do not cause them. To really close them completely you would need to avoid sharing anything, or not allowing to measure time, neither of which is practical short of an air gap. There are reasonable assessments you can make either way and the answers will be different based on your requirements. There isn't a single answer that works for everyone. There are cases where it isn't a problem at all. If you don't have multiple users on the system your tolerances should be extremely high. For users who have multiple users there can be different tradeoffs. So there isn't a single answer, and that is why it is important that this is configurable. -Andi"

Label: 1 (uncivil)

Email: b"O On Tue, 30 Jan 2018 17:14:45 +0200 Igor Stoppa <igor.stoppa@huawei.com> wrote: Please don't put plain-text files into core-api - that's a directory full of RST documents. Your document is 99.9% RST already, better to just finish the job and tie it into the rest of the kernel docs. We might as well put the SPDX tag here, it's a new file. This is all good information, but I'd suggest it belongs more in the 0/npatch posting than here. The introduction of *this* document should say what it actually covers. This seems like a relevant and important aspect of the API that shouldn't be buried in the middle of a section talking about random things. So one gets this far, but has no actual idea of how to do these things. Which leads me to wonder: what is this document for? Who are you expecting to read it? You could improve things a lot by (once again) going to RST and using directives to bring in the kernel doc comments from the source (which, Inote, do exist). But I'd suggest rethinking this document and its audience. Most of the people reading it are likely wanting to learn how to *use* this API; I think it would be best to not leave them frustrated. Thanks, jon"

Label: 1 (uncivil)

```
In [12]: bert_model_name = 'small_bert/bert_en_uncased_L-4_H-512_A-8'

map_name_to_handle = {
    'bert_en_uncased_L-12_H-768_A-12':
        'https://tfhub.dev/tensorflow/bert_en_uncased_L-12_H-768_A-12/3',
    'bert_en_cased_L-12_H-768_A-12':
        'https://tfhub.dev/tensorflow/bert_en_cased_L-12_H-768_A-12/3',
    'bert_multi_cased_L-12_H-768_A-12':
        'https://tfhub.dev/tensorflow/bert_multi_cased_L-12_H-768_A-12/3',
    'small_bert/bert_en_uncased_L-2_H-128_A-2':
        'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-2_H-128_A-2/1',
    'small_bert/bert_en_uncased_L-2_H-256_A-4':
        'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-2_H-256_A-4/1',
    'small_bert/bert_en_uncased_L-2_H-512_A-8':
        'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-2_H-512_A-8/1',
    'small_bert/bert_en_uncased_L-2_H-768_A-12':
        'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-2_H-768_A-12/1',
    'small_bert/bert_en_uncased_L-4_H-128_A-2':
        'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-4_H-128_A-2/1',
    'small_bert/bert_en_uncased_L-4_H-256_A-4':
        'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-4_H-256_A-4/1',
    'small_bert/bert_en_uncased_L-4_H-512_A-8':
        'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-4_H-512_A-8/1',
    'small_bert/bert_en_uncased_L-4_H-768_A-12':
        'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-4_H-768_A-12/1',
    'small_bert/bert_en_uncased_L-6_H-128_A-2':
        'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-6_H-128_A-2/1',
    'small_bert/bert_en_uncased_L-6_H-256_A-4':
        'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-6_H-256_A-4/1',
    'small_bert/bert_en_uncased_L-6_H-512_A-8':
        'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-6_H-512_A-8/1',
    'small_bert/bert_en_uncased_L-6_H-768_A-12':
        'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-6_H-768_A-12/1',
    'small_bert/bert_en_uncased_L-8_H-128_A-2':
        'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-8_H-128_A-2/1',
    'small_bert/bert_en_uncased_L-8_H-256_A-4':
        'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-8_H-256_A-4/1',
    'small_bert/bert_en_uncased_L-8_H-512_A-8':
        'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-8_H-512_A-8/1',
    'small_bert/bert_en_uncased_L-8_H-768_A-12':
        'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-8_H-768_A-12/1'
}
```

```

    'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-8_H-512_A-8/1',
    'small_bert/bert_en_uncased_L-8_H-768_A-12':
    'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-8_H-768_A-12/1',
    'small_bert/bert_en_uncased_L-10_H-128_A-2':
    'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-10_H-128_A-2/1',
    'small_bert/bert_en_uncased_L-10_H-256_A-4':
    'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-10_H-256_A-4/1',
    'small_bert/bert_en_uncased_L-10_H-512_A-8':
    'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-10_H-512_A-8/1',
    'small_bert/bert_en_uncased_L-10_H-768_A-12':
    'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-10_H-768_A-12/1',
    'small_bert/bert_en_uncased_L-12_H-128_A-2':
    'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-12_H-128_A-2/1',
    'small_bert/bert_en_uncased_L-12_H-256_A-4':
    'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-12_H-256_A-4/1',
    'small_bert/bert_en_uncased_L-12_H-512_A-8':
    'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-12_H-512_A-8/1',
    'small_bert/bert_en_uncased_L-12_H-768_A-12':
    'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-12_H-768_A-12/1',
    'albert_en_base':
    'https://tfhub.dev/tensorflow/albert_en_base/2',
    'electra_small':
    'https://tfhub.dev/google/electra_small/2',
    'electra_base':
    'https://tfhub.dev/google/electra_base/2',
    'experts_pubmed':
    'https://tfhub.dev/google/experts/bert/pubmed/2',
    'experts_wiki_books':
    'https://tfhub.dev/google/experts/bert/wiki_books/2',
    'talking-heads_base':
    'https://tfhub.dev/tensorflow/talkheads_ggelu_bert_en_base/1',
}

map_model_to_preprocess = {
    'bert_en_uncased_L-12_H-768_A-12':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'bert_en_cased_L-12_H-768_A-12':
    'https://tfhub.dev/tensorflow/bert_en_cased_preprocess/3',
    'small_bert/bert_en_uncased_L-2_H-128_A-2':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-2_H-256_A-4':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-2_H-512_A-8':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-2_H-768_A-12':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-4_H-128_A-2':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-4_H-256_A-4':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-4_H-512_A-8':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-4_H-768_A-12':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-6_H-128_A-2':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-6_H-256_A-4':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-6_H-512_A-8':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-6_H-768_A-12':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-8_H-128_A-2':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
    'small_bert/bert_en_uncased_L-8_H-256_A-4':

```

```

'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-8_H-512_A-8':
'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-8_H-768_A-12':
'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-10_H-128_A-2':
'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-10_H-256_A-4':
'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-10_H-512_A-8':
'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-10_H-768_A-12':
'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-12_H-128_A-2':
'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-12_H-256_A-4':
'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-12_H-512_A-8':
'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-12_H-768_A-12':
'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'bert_multi_cased_L-12_H-768_A-12':
'https://tfhub.dev/tensorflow/bert_multi_cased_preprocess/3',
'albert_en_base':
'https://tfhub.dev/tensorflow/albert_en_preprocess/3',
'electra_small':
'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'electra_base':
'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'experts_pubmed':
'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'experts_wiki_books':
'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'talking-heads_base':
'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
}

```

```

tfhub_handle_encoder = map_name_to_handle[bert_model_name]
tfhub_handle_preprocess = map_model_to_preprocess[bert_model_name]

```

```

print(f'BERT model selected          : {tfhub_handle_encoder}')
print(f'Preprocess model auto-selected: {tfhub_handle_preprocess}')

```

```

BERT model selected          : https://tfhub.dev/tensorflow/small_bert/bert_en_unca
sed_L-4_H-512_A-8/1
Preprocess model auto-selected: https://tfhub.dev/tensorflow/bert_en_uncased_preproc
ess/3

```

Preprocessing model

```
In [13]: bert_preprocess_model = hub.KerasLayer(tfhub_handle_preprocess)
```

```
In [14]: text_test = ["The driver is looking good!\n\nIt looks like you've done some kind of
text_preprocessed = bert_preprocess_model(text_test)
```

```

print(f'Keys          : {list(text_preprocessed.keys())}')
print(f'Shape         : {text_preprocessed["input_word_ids"].shape}')
print(f'Word Ids       : {text_preprocessed["input_word_ids"][0, :12]}')
print(f'Input Mask     : {text_preprocessed["input_mask"][0, :12]}')
print(f'Type Ids       : {text_preprocessed["input_type_ids"][0, :12]}')

```

```

Keys          : ['input_word_ids', 'input_mask', 'input_type_ids']
Shape         : (1, 128)
Word Ids      : [ 101 1996 4062 2003 2559 2204  999 2009 3504 2066 2017 1005]
Input Mask    : [1 1 1 1 1 1 1 1 1 1 1 1]
Type Ids      : [0 0 0 0 0 0 0 0 0 0 0 0]

```

Using BERT model

```
In [15]: bert_model = hub.KerasLayer(tfhub_handle_encoder)
```

```
In [16]: bert_results = bert_model(text_preprocessed)
```

```
print(f'Loaded BERT: {tfhub_handle_encoder}')
print(f'Pooled Outputs Shape:{bert_results["pooled_output"].shape}')
print(f'Pooled Outputs Values:{bert_results["pooled_output"][0, :12]}')
print(f'Sequence Outputs Shape:{bert_results["sequence_output"].shape}')
print(f'Sequence Outputs Values:{bert_results["sequence_output"][0, :12]}')
```

```
Loaded BERT: https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-4_H-512_A-8/1
Pooled Outputs Shape:(1, 512)
Pooled Outputs Values:[ 0.84747237  0.9954414  -0.28012952  0.12758777  0.31347865
 0.9054933
 0.51660377 -0.99680704 -0.0568257  -0.99887776  0.14184111 -0.98870677]
Sequence Outputs Shape:(1, 128, 512)
Sequence Outputs Values:[[ 0.39939746 -0.39085412  0.93853116 ... 0.2800356  0.033
86158
 -0.40618896]
 [-0.29228687  0.40331316 -1.0200574  ... -0.5753818  0.06500214
 0.86555773]
 [-0.8361566  0.07805394  0.6440221  ... 0.61097324  0.5496331
 0.5941888 ]
 ...
 [-0.31816947 -1.1716307  -1.4007772  ... 0.5933535  -0.5400041
 -0.59103024]
 [-0.40100175  0.1862402  -0.27396035  ... 0.6435042  0.38049674
 0.53075415]
 [-0.17033997  0.2594924  0.6192257  ... -0.47313097  0.6680391
 0.01982243]]
```

Defining model

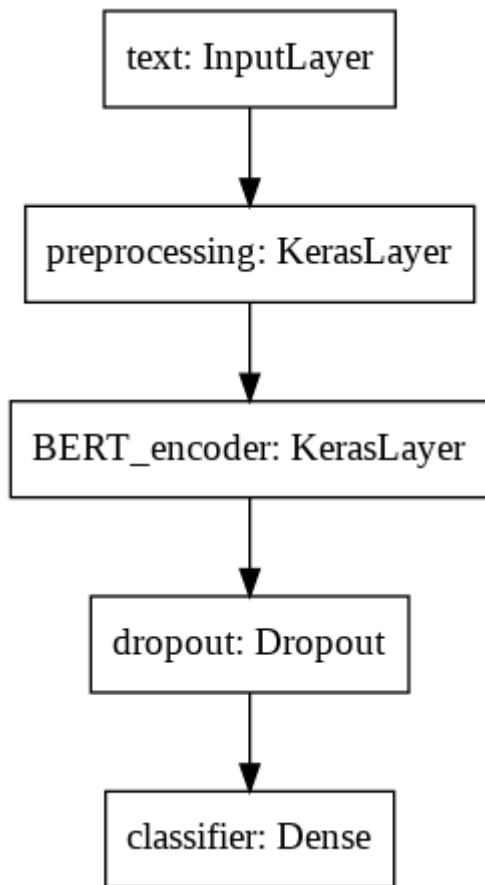
```
In [17]: def build_classifier_model():
text_input = tf.keras.layers.Input(shape=(), dtype=tf.string, name='text')
preprocessing_layer = hub.KerasLayer(tfhub_handle_preprocess, name='preprocessing')
encoder_inputs = preprocessing_layer(text_input)
encoder = hub.KerasLayer(tfhub_handle_encoder, trainable=True, name='BERT_encoder')
outputs = encoder(encoder_inputs)
net = outputs['pooled_output']
net = tf.keras.layers.Dropout(0.1)(net)
net = tf.keras.layers.Dense(1, activation=None, name='classifier')(net)
return tf.keras.Model(text_input, net)
```

```
In [18]: classifier_model = build_classifier_model()
bert_raw_result = classifier_model(tf.constant(text_test))
print(tf.sigmoid(bert_raw_result))
```

```
tf.Tensor([[0.6869482]], shape=(1, 1), dtype=float32)
```

```
In [19]: tf.keras.utils.plot_model(classifier_model)
```

```
Out[19]:
```

Model training:

```
In [20]: loss = tf.keras.losses.BinaryCrossentropy(from_logits=True)
        metrics = tf.metrics.BinaryAccuracy()
```

```
In [21]: epochs = 5
        steps_per_epoch = tf.data.experimental.cardinality(train_ds).numpy()
        num_train_steps = steps_per_epoch * epochs
        num_warmup_steps = int(0.1*num_train_steps)

        init_lr = 3e-5
        optimizer = optimization.create_optimizer(init_lr=init_lr,
                                                  num_train_steps=num_train_steps,
                                                  num_warmup_steps=num_warmup_steps,
                                                  optimizer_type='adamw')
```

```
In [22]: classifier_model.compile(optimizer=optimizer,
                                loss=loss,
                                metrics=metrics)
```

```
In [23]: print(f'Training model with {tfhub_handle_encoder}')
        history = classifier_model.fit(x=train_ds,
                                       validation_data=val_ds,
                                       epochs=epochs)
```

```
Training model with https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-4_H-51
2_A-8/1
Epoch 1/5
4/4 [=====] - 40s 8s/step - loss: 0.6847 - binary_accuracy:
0.5833 - val_loss: 0.6185 - val_binary_accuracy: 0.6296
Epoch 2/5
4/4 [=====] - 33s 8s/step - loss: 0.6370 - binary_accuracy:
0.5556 - val_loss: 0.6272 - val_binary_accuracy: 0.5556
Epoch 3/5
4/4 [=====] - 33s 8s/step - loss: 0.5698 - binary_accuracy:
0.7037 - val_loss: 0.6109 - val_binary_accuracy: 0.6667
```


Epoch 4/5
 4/4 [=====] - 32s 8s/step - loss: 0.5292 - binary_accuracy: 0.7315 - val_loss: 0.5996 - val_binary_accuracy: 0.6667
 Epoch 5/5
 4/4 [=====] - 33s 8s/step - loss: 0.4939 - binary_accuracy: 0.8056 - val_loss: 0.5961 - val_binary_accuracy: 0.6667

Evaluating:

```
In [24]: loss, accuracy = classifier_model.evaluate(test_ds)

print(f'Loss: {loss}')
print(f'Accuracy: {accuracy}')
```

2/2 [=====] - 3s 109ms/step - loss: 0.5933 - binary_accuracy: 0.6667
 Loss: 0.5932566523551941
 Accuracy: 0.6666666865348816

```
In [25]: history_dict = history.history
print(history_dict.keys())

acc = history_dict['binary_accuracy']
val_acc = history_dict['val_binary_accuracy']
loss = history_dict['loss']
val_loss = history_dict['val_loss']


epochs = range(1, len(acc) + 1)
fig = plt.figure(figsize=(10, 6))
fig.tight_layout()

plt.subplot(2, 1, 1)
# "bo" is for "blue dot"
plt.plot(epochs, loss, 'r', label='Training loss')
# b is for "solid blue line"
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
# plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.subplot(2, 1, 2)
plt.plot(epochs, acc, 'r', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')
```

```
dict_keys(['loss', 'binary_accuracy', 'val_loss', 'val_binary_accuracy'])
```

```
Out[25]: <matplotlib.legend.Legend at 0x7fcb136f2450>
```



```
In [26]: # testing:
examples = ["I would really like to get an ack from the people who have been \
deep into this first. If you can get that, and preferably resubmit with a \
less condescending changelog, I can pick it up.", "I like the idea and I think \
it's good direction to go, but could you please share some from perf stat or \
whatever you used to measure the new performance?", "What's advertisement \
there? Huch? Care to tell what's a lie instead of making bold statements? \
Thanks, tglx", "There probably is a decent compromise to find between \
'not accepting a single additional byte' and accepting several GB. \
For example how likely is it that the growth of this structure make it \
go over a page? I would hope not at all. By choosing a large but decent \
high limit, I think we can find a future-compatible compromise that doesn't \
rely on a preliminary getsockopt() just for structure truncation decision..."]

def print_results(inputs, results):
    for i in range(len(inputs)):
        prediction = "Uncivil"
        if results[i][0] >= 0.5:
            prediction = "Civil"
        print("Input:", inputs[i], "\nScore:", results[i][0], "\nPrediction:", prediction)

results = tf.sigmoid(classifier_model(tf.constant(examples)))
print_results(examples, results)
```

Input: I would really like to get an ack from the people who have been deep into this first. If you can get that, and preferably resubmit with a less condescending changelog, I can pick it up.

Score: tf.Tensor(0.67575943, shape=(), dtype=float32)

Prediction: Civil

Input: I like the idea and I think it's good direction to go, but could you please share some from perf stat or whatever you used to measure the new performance?

Score: tf.Tensor(0.6564582, shape=(), dtype=float32)

Prediction: Civil

Input: What's advertisement there? Huch? Care to tell what's a lie instead of making bold statements? Thanks, tglx

Score: tf.Tensor(0.5297446, shape=(), dtype=float32)

Prediction: Civil

Input: There probably is a decent compromise to find between 'not accepting a single additional byte' and accepting several GB. For example how likely is it that the growth of this structure make it go over a page? I would hope not at all. By choosing a large but decent high limit, I think we can find a future-compatible compromise that

doesn't rely on a preliminary getsockopt() just for structure truncation decision...
Score: tf.Tensor(0.7734759, shape=(), dtype=float32)
Prediction: Civil