

```
In [1]: # A dependency of the preprocessing for BERT inputs
!pip install -q tensorflow-text
```

```
4.3MB 6.3MB/s
454.3MB 34kB/s
4.0MB 13.4MB/s
6.0MB 45.4MB/s
1.2MB 29.2MB/s
4.0MB 25.8MB/s
471kB 56.1MB/s
4.9MB 41.3MB/s
```

```
In [2]: # Using AdamW optimizer
!pip install -q tf-models-official==2.4
```

```
1.1MB 5.2MB/s
645kB 18.2MB/s
686kB 21.0MB/s
38.2MB 76kB/s
358kB 46.1MB/s
102kB 14.7MB/s
174kB 54.9MB/s
51kB 8.8MB/s
1.2MB 44.1MB/s
```

Building wheel for py-cpuinfo (setup.py) ... done

Building wheel for sequeval (setup.py) ... done

```
In [3]: import os
import shutil
import tensorflow as tf
import tensorflow_hub as hub
import tensorflow_text as text
from official.nlp import optimization # to create AdamW optimizer
import matplotlib.pyplot as plt
tf.get_logger().setLevel('ERROR')
```

```
In [4]: url = 'https://github.com/ahlralf/point/blob/main/v2_emails.tar.gz?raw=true'
dataset = tf.keras.utils.get_file('v2_emails.tar.gz', url,
                                untar=True, cache_dir='.',
                                cache_subdir='')
```

Downloading data from https://github.com/ahlralf/point/blob/main/v2_emails.tar.gz?raw=true

434176/429289 [=====] - 0s 0us/step

```
In [5]: AUTOTUNE = tf.data.AUTOTUNE
batch_size = 32
seed = 42

raw_train_ds = tf.keras.preprocessing.text_dataset_from_directory(
    'v2_emails/train',
    batch_size=batch_size,
    validation_split=0.2,
    subset='training',
    seed=seed)

class_names = raw_train_ds.class_names
train_ds = raw_train_ds.cache().prefetch(buffer_size=AUTOTUNE)

val_ds = tf.keras.preprocessing.text_dataset_from_directory(
    'v2_emails/train',
    batch_size=batch_size,
    validation_split=0.2,
    subset='validation',
    seed=seed)
```

```

val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)

test_ds = tf.keras.preprocessing.text_dataset_from_directory(
    'v2_emails/test',
    batch_size=batch_size)

test_ds = test_ds.cache().prefetch(buffer_size=AUTOTUNE)

```

Found 1238 files belonging to 2 classes.
 Using 991 files for training.
 Found 1238 files belonging to 2 classes.
 Using 247 files for validation.
 Found 308 files belonging to 2 classes.

Preprocessing and looking at a few emails:

```

In [6]: import re
import nltk
nltk.download('stopwords')
from nltk.corpus import stopwords

```

[nltk_data] Downloading package stopwords to /root/nltk_data...
 [nltk_data] Unzipping corpora/stopwords.zip.

```

In [7]: regex_tokenizer = nltk.RegexpTokenizer("\w+")
def text_preprocessing(content):
    content = str(content)
    content = re.sub("[^a-zA-Z]", " ", content)
    content = content.lower()
    content = content.encode("utf-8", "ignore").decode()
    content = " ".join(regex_tokenizer.tokenize(content))
    for c in content:
        c.replace('\n', ' ')
    words = content.split()
    stops = set(stopwords.words("english"))
    words = [w for w in words if not w in stops]

    return ' '.join(words)

train_2 = train_ds
for text_batch, label_batch in train_2:
    text_batch = text_preprocessing(text_batch)

for text_batch, label_batch in train_2.take(1):
    for i in range(10):
        print(f'Email: {text_batch.numpy()[i]}')
        label = label_batch.numpy()[i]
        print(f'Label: {label} ({class_names[label]})')

```

Email: b"Hi Ulf, Previous multi slot implementation was removed as nobody used it and nobody tested it. There are lots of mistakes in previous implementation which are not related to request serialization like lack of slot switch / lack of adding slot id to CIU commands / etc... So obviously it was never tested and never used at real multi slot hardware. In current implementation data transfers and commands to different hosts (slots) are serialized internally in the dw_mmc driver. We have request queue and when .request() is called we add new request to the queue. We take new request from the queue only if the previous one has already finished. So although hosts (slots) have separate locks (mmc_claim|release_host()) the requests to different slots are serialized by driver. Isn't that enough? I'm not very familiar with SD/SDIO/(e)MMC specs so my assumptions might be wrong in that case please correct me. Nevertheless we had to deal somehow with existing hardware which has multi slot dw_mmc controller and both slots are used... This patch at least shouldn't break anything for current users (which use it in single slot mode). Moreover we tested this dual-slot implementation and don't catch any problems (probably yet) except bus performance decrease in dual-slot mode (which is quite expected). -- Eugeniy Paltsev"

Label: 1 (technical)

Email: b"Hi, Eduardo, as it is late in this merge window, I'd prefer to 1. drop all the

e thermal-soc material in the first pull request which I will send out soon.2. you can prepare another pull request containing the thermal-soc materials except the exynos fixes3. exynos fixes with the problem solved can be queued for -rc2 or later. thanks, rui"

Label: 1 (technical)

Email: b'Sure, it leaves the function to deal with the equiv table length only and the caller then adds the header length. Which is actually cleaner.--Regards/Gruss, Boris. Good mailing practices for 400: avoid top-posting and trim the reply.'

Label: 1 (technical)

Email: b"I'm not sure I like this. If you have a userspace application built against more recent uapi headers than the kernel you are actually running on, then by definition you won't have this check in place, and you'll get EINVAL returns anyway. If you just backport this patch to an older kernel, you'll not get the EINVAL return, but you will get silent failures on event subscriptions that your application thinks exist, but the kernel doesn't recognize. This would make sense if you had a way to communicate back to user space the unrecognized options, but since we don't (currently) have that, I would rather see the EINVAL returned than just have things not work. Neil"

Label: 1 (technical)

Email: b'Hi, Yes that is fine by me and you\'ve my permission to switch to using just the SPDX header. FWIW I do not believe the "can\'t be removed from \'this software and associated documentation files (the "Software")\'" language applies to the software as a whole and not individual files. Yes you may make the same change to all files with my copyright. Regards, Hans'

Label: 1 (technical)

Email: b"The bug fix is to handle non-vmalloc pages. I'll see if I can do a smaller and more bandaidd-y fix first."

Label: 1 (technical)

Email: b"I can't take patches without any changelog text at all :("

Label: 0 (nontechnical)

Email: b"Hi!Yes.No, I don't want that.I'd like 2 child nodes, each specifying which HVLEds it controls.Let me edit the original proposal.\t\t\t\t\t\t\t\t\t\tPavel--(english) <http://www.livejournal.com/~pavelmachek>(cesky, pictures) <http://atrey.karlin.mff.cuni.cz/~pavel/picture/horses/blog.html>"

Label: 1 (technical)

Email: b" I think the `__KERNEL__` and `asm/errno.h` slip-ups are things I cargo-culted from the arch code as a fresh-faced noob yet to learn the finer details, so ack for those parts. The forward-declarations, though, were a deliberate effort to minimise header dependencies and compilation bloat for inclusions who absolutely wouldn't care, and specifically to try to avoid setting transitive include expectations since they always seem to end up breaking someone's config somewhere down the line. Admittedly this little backwater is hardly comparable to the likes of the `sched.h` business, but I'm still somewhat on the fence about that change :/Robin."

Label: 1 (technical)

Email: b'Interesting \xe2\x80\xa6Would you like to share any more information from t
his meeting?I would appreciate further indications for a corresponding change accept
ance.I found a feedback by Mauro Carvalho Chehab more constructive.[GIT,PULL,FOR,v4.
15] Cleanup fixeshttps://patchwork.linuxtv.org/patch/43957/\xe2\x80\x9c\xe2\x80\xa6T
his time, I was nice and I took some time doing:\t\$ quilt fold < `quilt next` && qui
lt delete `quilt next`\xe2\x80\xa6\xe2\x80Regards,Markus'

Label: 0 (nontechnical)

Choosing BERT model to fine-tune (BERT-small)

```
In [8]: bert_model_name = 'small_bert/bert_en_uncased_L-4_H-512_A-8'

map_name_to_handle = {
    'bert_en_uncased_L-12_H-768_A-12':
        'https://tfhub.dev/tensorflow/bert_en_uncased_L-12_H-768_A-12/3',
    'bert_en_cased_L-12_H-768_A-12':
        'https://tfhub.dev/tensorflow/bert_en_cased_L-12_H-768_A-12/3',
    'bert_multi_cased_L-12_H-768_A-12':
        'https://tfhub.dev/tensorflow/bert_multi_cased_L-12_H-768_A-12/3',
    'small_bert/bert_en_uncased_L-2_H-128_A-2':
        'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-2_H-128_A-2/1',
    'small_bert/bert_en_uncased_L-2_H-256_A-4':
        'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-2_H-256_A-4/1',
    'small_bert/bert_en_uncased_L-2_H-512_A-8':
        'https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-2_H-512_A-8/1',
}
```

}

```

'small_bert/bert_en_uncased_L-2_H-512_A-8':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-2_H-768_A-12':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-4_H-128_A-2':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-4_H-256_A-4':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-4_H-512_A-8':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-4_H-768_A-12':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-6_H-128_A-2':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-6_H-256_A-4':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-6_H-512_A-8':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-6_H-768_A-12':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-8_H-128_A-2':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-8_H-256_A-4':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-8_H-512_A-8':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-8_H-768_A-12':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-10_H-128_A-2':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-10_H-256_A-4':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-10_H-512_A-8':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-10_H-768_A-12':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-12_H-128_A-2':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-12_H-256_A-4':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-12_H-512_A-8':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'small_bert/bert_en_uncased_L-12_H-768_A-12':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'bert_multi_cased_L-12_H-768_A-12':
    'https://tfhub.dev/tensorflow/bert_multi_cased_preprocess/3',
'albert_en_base':
    'https://tfhub.dev/tensorflow/albert_en_preprocess/3',
'electra_small':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'electra_base':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'experts_pubmed':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'experts_wiki_books':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
'talking-heads_base':
    'https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3',
}

tfhub_handle_encoder = map_name_to_handle[bert_model_name]
tfhub_handle_preprocess = map_model_to_preprocess[bert_model_name]

print(f'BERT model selected           : {tfhub_handle_encoder}')
print(f'Preprocess model auto-selected: {tfhub_handle_preprocess}')

```

BERT model selected : https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-4_H-512_A-8/1
 Preprocess model auto-selected: https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3

Preprocessing model

```
In [9]: bert_preprocess_model = hub.KerasLayer(tfhub_handle_preprocess)
```

```
In [10]: text_test = ["The driver is looking good!It looks like you've done some kind of revi
text_preprocessed = bert_preprocess_model(text_test)

print(f'Keys      : {list(text_preprocessed.keys())}')
print(f'Shape     : {text_preprocessed["input_word_ids"].shape}')
print(f'Word Ids   : {text_preprocessed["input_word_ids"][0, :12]}')
print(f'Input Mask : {text_preprocessed["input_mask"][0, :12]}')
print(f'Type Ids   : {text_preprocessed["input_type_ids"][0, :12]}')

Keys      : ['input_word_ids', 'input_mask', 'input_type_ids']
Shape     : (1, 128)
Word Ids   : [ 101 1996 4062 2003 2559 2204  999 2009 3504 2066 2017 1005]
Input Mask : [1 1 1 1 1 1 1 1 1 1 1 1]
Type Ids   : [0 0 0 0 0 0 0 0 0 0 0 0]
```

Using BERT model

```
In [11]: bert_model = hub.KerasLayer(tfhub_handle_encoder)
```

```
In [12]: bert_results = bert_model(text_preprocessed)

print(f'Loaded BERT: {tfhub_handle_encoder}')
print(f'Pooled Outputs Shape:{bert_results["pooled_output"].shape}')
print(f'Pooled Outputs Values:{bert_results["pooled_output"][0, :12]}')
print(f'Sequence Outputs Shape:{bert_results["sequence_output"].shape}')
print(f'Sequence Outputs Values:{bert_results["sequence_output"][0, :12]}')

Loaded BERT: https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-4_H-512_A-8/1
Pooled Outputs Shape:(1, 512)
Pooled Outputs Values:[ 0.84747237  0.9954414  -0.28012952  0.12758777  0.31347865
 0.9054933
 0.51660377 -0.99680704 -0.0568257  -0.99887776  0.14184111 -0.98870677]
Sequence Outputs Shape:(1, 128, 512)
Sequence Outputs Values:[[ 0.39939746 -0.39085412  0.93853116 ... 0.2800356  0.033
86158
 -0.40618896]
 [-0.29228687  0.40331316 -1.0200574  ... -0.5753818  0.06500214
 0.86555773]
 [-0.8361566  0.07805394  0.6440221  ... 0.61097324  0.5496331
 0.5941888 ]
 ...
 [-0.31816947 -1.1716307  -1.4007772  ... 0.5933535  -0.5400041
 -0.59103024]
 [-0.40100175  0.1862402  -0.27396035 ... 0.6435042  0.38049674
 0.53075415]
 [-0.17033997  0.2594924  0.6192257  ... -0.47313097  0.6680391
 0.01982243]]
```

The BERT models return a map with 3 important keys: pooled_output, sequence_output, encoder_outputs:

"pooled_output" represents each input sequence as a whole. The shape is [batch_size, H]. [~ Embedding for the entire email]
 "sequence_output" represents each input token in the context. The shape is [batch_size, seq_length, H]. [~ contextual embedding for every token in the email]
 "encoder_outputs" are the intermediate activations of the L Transformer blocks.

outputs["encoder_outputs"][i] is a Tensor of shape [batch_size, seq_length, 1024] with the outputs of the i-th Transformer block, for $0 \leq i < L$. The last value of the list is equal to sequence_output.

For the fine-tuning we use the pooled_output array.

Defining the model

Fine-tuned model comprising preprocessing model + selected BERT model + 1 dense + 1 dropout layer

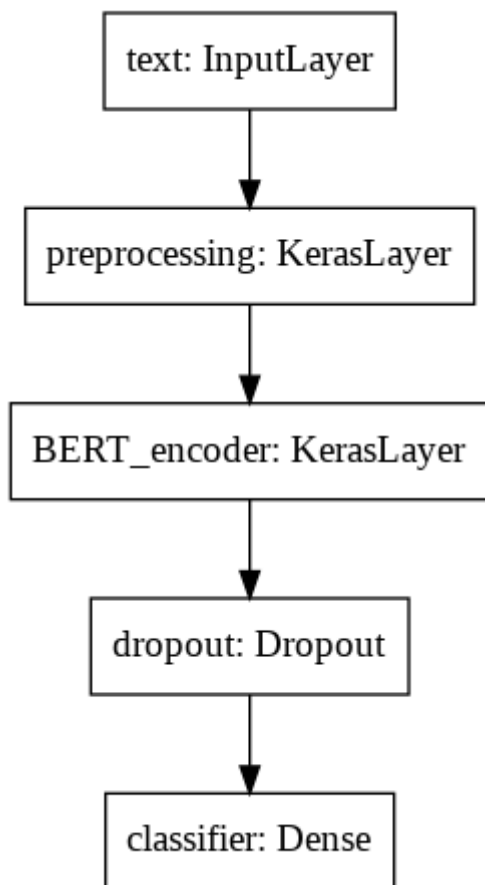
```
In [14]: def build_classifier_model():
text_input = tf.keras.layers.Input(shape=(), dtype=tf.string, name='text')
preprocessing_layer = hub.KerasLayer(tfhub_handle_preprocess, name='preprocessing')
encoder_inputs = preprocessing_layer(text_input)
encoder = hub.KerasLayer(tfhub_handle_encoder, trainable=True, name='BERT_encoder')
outputs = encoder(encoder_inputs)
net = outputs['pooled_output']
net = tf.keras.layers.Dropout(0.1)(net)
net = tf.keras.layers.Dense(1, activation=None, name='classifier')(net)
return tf.keras.Model(text_input, net)
```

```
In [15]: classifier_model = build_classifier_model()
bert_raw_result = classifier_model(tf.constant(text_test))
print(tf.sigmoid(bert_raw_result))
```

tf.Tensor([[0.5468524]], shape=(1, 1), dtype=float32)

```
In [16]: tf.keras.utils.plot_model(classifier_model)
```

Out[16]:



Model training

Loss function for binary classification


```
In [17]: loss = tf.keras.losses.BinaryCrossentropy(from_logits=True)
metrics = tf.metrics.BinaryAccuracy()
```

Optimizer: AdamW

For the learning rate (init_lr), we use the same schedule as BERT pre-training: linear decay of a notional initial learning rate, prefixed with a linear warm-up phase over the first 10% of training steps (num_warmup_steps). In line with the BERT paper, the initial learning rate is smaller for fine-tuning (best of 5e-5, 3e-5, 2e-5).

```
In [18]: epochs = 5
steps_per_epoch = tf.data.experimental.cardinality(train_ds).numpy()
num_train_steps = steps_per_epoch * epochs
num_warmup_steps = int(0.1*num_train_steps)

init_lr = 3e-5
optimizer = optimization.create_optimizer(init_lr=init_lr,
                                         num_train_steps=num_train_steps,
                                         num_warmup_steps=num_warmup_steps,
                                         optimizer_type='adamw')
```

Loading BERT model and training

```
In [19]: classifier_model.compile(optimizer=optimizer,
                                loss=loss,
                                metrics=metrics)
```

```
In [20]: print(f'Training model with {tfhub_handle_encoder}')
history = classifier_model.fit(x=train_ds,
                              validation_data=val_ds,
                              epochs=epochs)
```

Training model with https://tfhub.dev/tensorflow/small_bert/bert_en_uncased_L-4_H-512_A-8/1

Epoch 1/5

31/31 [=====] - 226s 7s/step - loss: 0.4785 - binary_accuracy: 0.7326 - val_loss: 0.3047 - val_binary_accuracy: 0.9150

Epoch 2/5

31/31 [=====] - 219s 7s/step - loss: 0.3326 - binary_accuracy: 0.8860 - val_loss: 0.2732 - val_binary_accuracy: 0.9271

Epoch 3/5

31/31 [=====] - 222s 7s/step - loss: 0.2998 - binary_accuracy: 0.8900 - val_loss: 0.2589 - val_binary_accuracy: 0.9069

Epoch 4/5

31/31 [=====] - 223s 7s/step - loss: 0.2772 - binary_accuracy: 0.8961 - val_loss: 0.2513 - val_binary_accuracy: 0.8907

Epoch 5/5

31/31 [=====] - 220s 7s/step - loss: 0.2505 - binary_accuracy: 0.9072 - val_loss: 0.2323 - val_binary_accuracy: 0.9393

Evaluating model:

```
In [21]: loss, accuracy = classifier_model.evaluate(test_ds)

print(f'Loss: {loss}')
print(f'Accuracy: {accuracy}')
```

10/10 [=====] - 18s 2s/step - loss: 0.2901 - binary_accuracy: 0.8961

Loss: 0.29012614488601685

Accuracy: 0.8961039185523987

Plotting accuracy, loss over time:

```
In [22]: history_dict = history.history
print(history_dict.keys())
```



```

acc = history_dict['binary_accuracy']
val_acc = history_dict['val_binary_accuracy']
loss = history_dict['loss']
val_loss = history_dict['val_loss']

epochs = range(1, len(acc) + 1)
fig = plt.figure(figsize=(10, 6))
fig.tight_layout()


plt.subplot(2, 1, 1)
# "bo" is for "blue dot"
plt.plot(epochs, loss, 'r', label='Training loss')
# b is for "solid blue line"
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
# plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.subplot(2, 1, 2)
plt.plot(epochs, acc, 'r', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')

```

```
dict_keys(['loss', 'binary_accuracy', 'val_loss', 'val_binary_accuracy'])
```

```
Out[22]: <matplotlib.legend.Legend at 0x7f2075c63f10>
```



```

In [23]: # testing:
examples = ["Was looking for that. Thanks. Speaking of that, recent lksctp-tools \
got some defines to help knowing which features the available kernel headers \
have as it now probes if specific struct members are available or not. \
Though yeah, it also wouldn't help in this case, just mentioning it.",
"Sorry but I don't like imposing a run-time check on everybody \
when stack-based requests are the odd ones out. If we're going to make \
this a run-time check (I'd much prefer a compile-time check, but I \
understand that this may involve too much churn), then please do it \

```

```

for stack-based request users only.", "And I was just reminded about huge \
pages. But still, my point of finding a compromise still stands.", \
"Since when is the cover letter \
mandatory? I understand that is helps for a complicated patch set \
to explain the problem and solution in the cover letter, but for this \
simple test case addition what's the point? And there is nothing \
forcing a cover letter in", "I'm not exactly sure how Linux switch driver \
works, but from DT perspective I think we should rather have \
*hardware* described instead of a common Linux case. If I'm right, \
we should rather have all 3 switch ports described (5, 7,8) and have \
Linux just use the one it needs." ]
# technical, non-technical, technical, non-technical, technical

def print_results(inputs, results):
    for i in range(len(inputs)):
        prediction = "Non-technical"
        if results[i][0]>=0.5:
            prediction = "Technical"
        print("Input:", inputs[i], "\nScore:", results[i][0], "\nPrediction:", prediction)

results = tf.sigmoid(classifier_model(tf.constant(examples)))
print_results(examples, results)

```

Input: Was looking for that. Thanks. Speaking of that, recent lksctp-tools got some defines to help knowing which features the available kernel headers have as it now probes if specific struct members are available or not. Though yeah, it also wouldn't help in this case, just mentioning it.

Score: tf.Tensor(0.8919314, shape=(), dtype=float32)

Prediction: Technical

Input: Sorry but I don't like imposing a run-time check on everybody when stack-based requests are the odd ones out. If we're going to make this a run-time check (I'd much prefer a compile-time check, but I understand that this may involve too much churn), then please do it for stack-based request users only.

Score: tf.Tensor(0.77189916, shape=(), dtype=float32)

Prediction: Technical

Input: And I was just reminded about huge pages. But still, my point of finding a compromise still stands.

Score: tf.Tensor(0.9071373, shape=(), dtype=float32)

Prediction: Technical

Input: Since when is the cover letter mandatory? I understand that is helps for a complicated patch set to explain the problem and solution in the cover letter, but for this simple test case addition what's the point? And there is nothing forcing a cover letter in

Score: tf.Tensor(0.8546599, shape=(), dtype=float32)

Prediction: Technical

Input: I'm not exactly sure how Linux switch driver works, but from DT perspective I think we should rather have *hardware* described instead of a common Linux case. If I'm right, we should rather have all 3 switch ports described (5, 7,8) and have Linux just use the one it needs.

Score: tf.Tensor(0.96465874, shape=(), dtype=float32)

Prediction: Technical